

[Python-Dev] More compact dictionaries with faster iteration

Raymond Hettinger [raymond.hettinger at gmail.com](mailto:raymond.hettinger@gmail.com)

Mon Dec 10 02:44:30 CET 2012

- Previous message: [\[Python-Dev\] hg annotate is broken on hg.python.org](#)
 - Next message: [\[Python-Dev\] More compact dictionaries with faster iteration](#)
 - Messages sorted by: [\[date \]](#) [\[thread \]](#) [\[subject \]](#) [\[author \]](#)
-

The current memory layout for dictionaries is unnecessarily inefficient. It has a sparse table of 24-byte entries containing the hash value, key pointer, and value pointer.

Instead, the 24-byte entries should be stored in a dense table referenced by a sparse table of indices.

For example, the dictionary:

```
d = {'timmy': 'red', 'barry': 'green', 'guido': 'blue'}
```

is currently stored as:

```
entries = [['--', '--', '--'],
           [-8522787127447073495, 'barry', 'green'],
           ['--', '--', '--'],
           ['--', '--', '--'],
           ['--', '--', '--'],
           [-9092791511155847987, 'timmy', 'red'],
           ['--', '--', '--'],
           [-6480567542315338377, 'guido', 'blue']]
```

Instead, the data should be organized as follows:

```
indices = [None, 1, None, None, None, 0, None, 2]
entries = [[-9092791511155847987, 'timmy', 'red'],
           [-8522787127447073495, 'barry', 'green'],
           [-6480567542315338377, 'guido', 'blue']]
```

Only the data layout needs to change. The hash table algorithms would stay the same. All of the current optimizations would be kept, including key-sharing dicts and custom lookup functions for string-only dicts. There is no change to the hash functions, the table search order, or collision statistics.

The memory savings are significant (from 30% to 95% compression depending on the how full the table is). Small dicts (size 0, 1, or 2) get the most benefit.

For a sparse table of size *t* with *n* entries, the sizes are:

```
curr_size = 24 * t
new_size = 24 * n + sizeof(index) * t
```

In the above timmy/barry/guido example, the current size is 192 bytes (eight 24-byte entries) and the new size is 80 bytes (three 24-byte entries plus eight 1-byte indices). That gives 58% compression.

Note, the `sizeof(index)` can be as small as a single byte for small dicts, two bytes for bigger dicts and up to `sizeof(Py_ssize_t)` for huge dict.

In addition to space savings, the new memory layout makes iteration faster. Currently, `keys()`, `values`, and `items()` loop over the sparse table, skipping-over free slots in the hash table. Now, `keys/values/items` can loop directly over the dense table, using fewer memory accesses.

Another benefit is that resizing is faster and touches fewer pieces of memory. Currently, every hash/key/value entry is moved or copied during a resize. In the new layout, only the indices are updated. For the most part, the hash/key/value entries never move (except for an occasional swap to fill a hole left by a deletion).

With the reduced memory footprint, we can also expect better cache utilization.

For those wanting to experiment with the design, there is a pure Python proof-of-concept here:

<http://code.activestate.com/recipes/578375>

YMMV: Keep in mind that the above size statics assume a build with 64-bit `Py_ssize_t` and 64-bit pointers. The space savings percentages are a bit different on other builds. Also, note that in many applications, the size of the data dominates the size of the container (i.e. the weight of a bucket of water is mostly the water, not the bucket).

Raymond

-
- Previous message: [\[Python-Dev\] hg annotate is broken on hg.python.org](#)
 - Next message: [\[Python-Dev\] More compact dictionaries with faster iteration](#)
 - **Messages sorted by:** [\[date\]](#) [\[thread\]](#) [\[subject\]](#) [\[author\]](#)
-

[More information about the Python-Dev mailing list](#)